

Linear Regression

Guowei Wei
Department of Mathematics
Michigan State University

References:

Duc D. Nguyen's lecture notes

Wikipedia

classification



clustering



category

predicting a
quantity

just
looking

predicting
structure

few features
should be

few features
should be
important



kernel approximation

dimensionality
reduction

Data sets

Labeled data sets for supervised learning:

Regression (R):

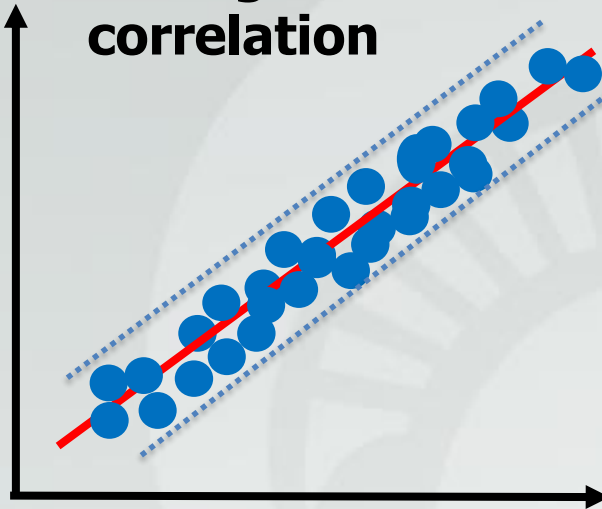
Data set (R): $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)}) \mid \mathbf{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \mathbb{R}\}_{i=1}^M$

Classification (C):


Data set (C): $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)}) \mid \mathbf{x}^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{-1, 1\}\}_{i=1}^M$

Correlations


**Strong
correlation**




**Weak
correlation**



**Strong
correlation**






**Non
correlation**



Linear Regression

In statistics, linear regression is a linear approach to modelling the relationship between a scalar response (or dependent variable) and one or more explanatory variables (or independent variables).




One Variable Linear Regression: Example

Assume we have a dataset giving the living areas and prices of **47** houses from Portland, Oregon:

Living area (feet ²)	Price (1000\$)
2104	400
1600	330
2400	369
1416	232
3000	540
⋮	⋮

One Variable Linear Regression: Example



Training/Test Sets

- In each house, we have living area (**feature**) and price (label)
- The previous dataset has given labels, thus we call it **training set**.
- If the dataset does **not** have labels, we call it **test set**

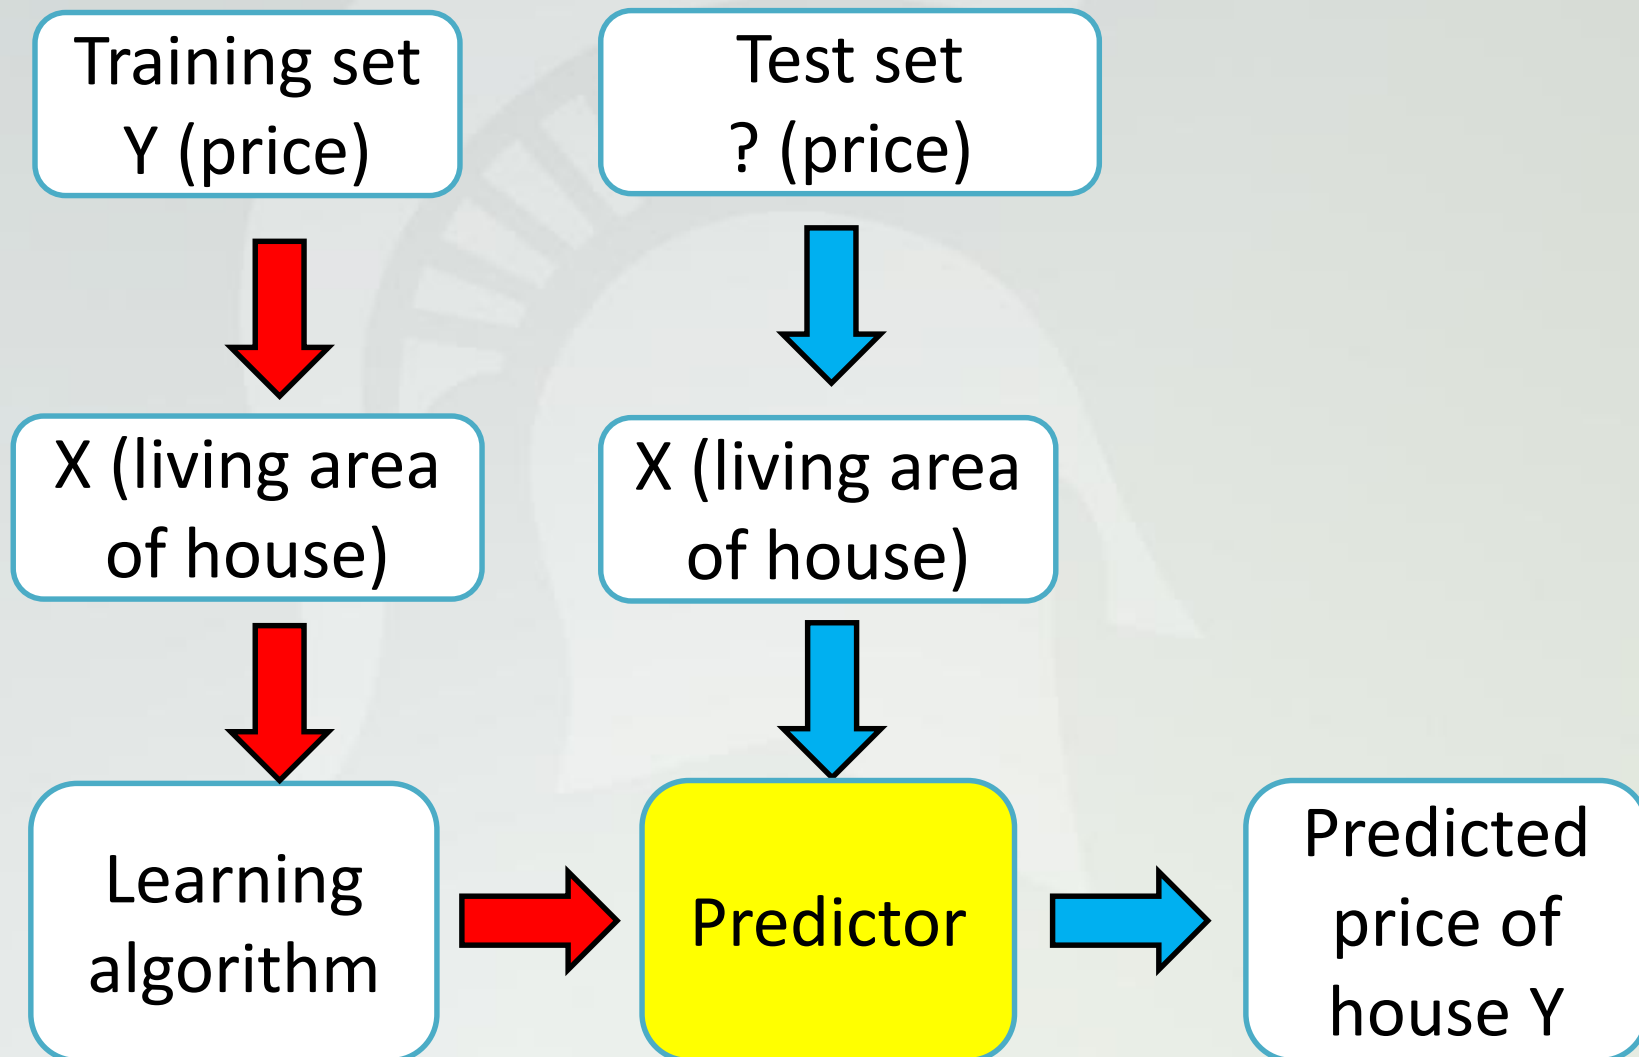
Living area (feet ²)	Price (1000\$s)
2104	400
1600	330
2400	369
1416	232
3000	540
⋮	⋮

Test set

- If we are given a size of living area in a house ,
What is the estimated price of that house?

Living area	Estimated Price
1300	?
4000	?
2200	?
2000	?

Model Representation



Predictor and Loss Function

- We assume a predictor that is linear in model parameter (c_0, c_1) :

$$p(x) = c_0 + c_1 x$$

- We choose c_0, c_1 such that they minimize the following **loss function**

$$L(c_0, c_1) = \sum_{i=1}^M (p(x^{(i)}) - y^{(i)})^2 = \|\mathbf{P} - \mathbf{Y}\|_2^2$$

where: $\mathbf{P} = (p(x^{(1)}), p(x^{(2)}), \dots, p(x^{(M)}))^T$

$$\mathbf{Y} = (y^{(1)}, y^{(2)}, \dots, y^{(M)})^T$$

Minimizing Loss Function

- In the dataset, $x^{(i)}$ and $y^{(i)}$ are, respectively, the living area and price of the i^{th} house. And $M = 45$

$$\min_{c_0, c_1} : L(c_0, c_1) = \sum_{i=1}^M (p(x^{(i)}) - y^{(i)})^2$$


is known as the **least-square linear regression problem**.

The optimal values of c_0, c_1 are:

$$\frac{\partial L}{\partial c_j} = 0, j = 0, 1 \Rightarrow$$

$$\hat{c}_1 = \frac{\sum_{i=1}^M x^{(i)} y^{(i)} - \frac{1}{M} \sum_{i=1}^M x^{(i)} \sum_{i=1}^M y^{(i)}}{\sum_{i=1}^M (x^{(i)})^2 - \frac{1}{M} \left(\sum_{i=1}^M x^{(i)} \right)^2}$$
$$\hat{c}_0 = \frac{1}{M} \sum_{i=1}^M y^{(i)} - \hat{c}_1 \frac{1}{M} \sum_{i=1}^M x^{(i)}$$

Result



L^p – norm


For real number $p \geq 1$, the L^p – norm of \mathbf{x} is:

$$\|\mathbf{x}\|_p = \left(\sum_{j=1}^n |x_j|^p \right)^{\frac{1}{p}}$$

The L^∞ – norm is:

$$\|\mathbf{x}\|_\infty = \max\{|x_1|, |x_2|, \dots, |x_n|\}$$

Figure: Illustration of L^p – norm. Every vector from the origin to the unit circle has a length of one.



[Image credit: Quartl](#)

Multiple Variables Linear Regression: Example

- Used when having multiple features
- In the housing example, consider a richer dataset with knowing the number of bedrooms in each house

x_1 Living area (feet ²)	x_2 #bedrooms	y Price (1000\$s)
2104	3	400
1600	3	330
2400	3	369
1416	2	232
3000	4	540
\vdots	\vdots	\vdots

Predictor and Loss Function

- We assume our predictor:

$$p(\mathbf{x}) = c_0 + c_1 x_1 + c_2 x_2$$

- Find c_0, c_1, c_2 to optimize the loss function:

$$L(c_0, c_1, c_2) = \sum_{i=1}^M \left(p(x_1^{(i)}, x_2^{(i)}) - y^{(i)} \right)^2$$

$$\frac{\partial L}{\partial c_j} = 0, j = 0, 1, 2 \Rightarrow$$

$$\begin{bmatrix} c_0 \\ c_1 \\ c_2 \end{bmatrix} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

Minimizing the Loss Function

- Solution of the optimization problem is $\begin{bmatrix} c_0 \\ c_1 \\ c_2 \end{bmatrix} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$

where $\mathbf{X} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} \\ \dots & \dots & \dots \\ 1 & x_1^{(M)} & x_2^{(M)} \end{bmatrix}$, and

$$\mathbf{Y} = (y^{(1)}, y^{(2)}, \dots, y^{(M)})$$

General linear regression model

- In general, we assume our predictor:

$$p(\mathbf{x}) = c_0 + c_1x_1 + \cdots + c_nx_n$$

Find c_0, c_1, \dots, c_n to optimize the loss function:

$$L(c_0, c_1, \dots, c_n) = \sum_{i=1}^M \left(p \left(x_1^{(i)}, \dots, x_n^{(i)} \right) \right)$$

General linear regression model

- Solution of the optimization problem is:




$$\begin{bmatrix} c_0 \\ c_1 \\ \dots \\ c_n \end{bmatrix} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

where $\mathbf{X} = \begin{bmatrix} 1 & x_1^{(1)} & \dots & x_n^{(1)} \\ 1 & x_1^{(2)} & \dots & x_n^{(2)} \\ \dots & \dots & \dots & \dots \\ 1 & x_1^{(M)} & \dots & x_n^{(M)} \end{bmatrix}$, and

$$\mathbf{Y} = (y^{(1)}, y^{(2)}, \dots, y^{(M)})$$

Discussions: Overfitting & linearity

- A model leads to overfitting when it perfectly fits the training data but poorly fits the test data



- Linear regression is about the linearity with respect to c not \mathbf{X}

Discussions: Loss Function minimization

- Least-square linear regression problem

$$\min_{c_0, c_1, \dots} : L(c_0, c_1, \dots) = \sum_{i=1}^M (p(\mathbf{x}^{(i)}) - y^{(i)})^2$$

- Gauss–Markov theorem: The above is the best linear unbiased estimator if the errors have expectation zero, are uncorrelated and have equal variances.
- Quantile regression: aims at estimating either the conditional median or other quantiles of the response variable
- Least absolute shrinkage and selection operator (**Lasso**)

Discussions: Loss Function minimization with L1 and L2 norms

$$L_1: \min_{c_0, c_1, \dots} : L(c_0, c_1, \dots) = \sum_{i=1}^M |p(\mathbf{x}^{(i)}) - y^{(i)}|$$

$$L_2: \min_{c_0, c_1, \dots} : L(c_0, c_1, \dots) = \sum_{i=1}^M (p(\mathbf{x}^{(i)}) - y^{(i)})^2$$

Least Squares Regression	Least Absolute Deviations Regression
Not very robust	Robust
Stable solution	Unstable solution
Always one solution	Possibly multiple solutions
No feature selection	Built-in feature selection
Non-sparse outputs	Sparse outputs
Computational efficient due to having analytical solutions	Computational inefficient on non-sparse cases